RL Tutorial 4

Marcelo Mattar

Learning vs. planning

- So far, we've seen how agents learn values by acting in the world and *experiencing* the outcome of their actions
- Methods such as Q-learning are ***model-free*** as they do not require a model

- In this tutorial, we will learn about ***model-based*** methods, which compute action values via planning.
- Instead of learning values from experience, planning is the process of computing action values from a model.

But what is a model?

A representation of how the world will respond to the agent's actions

But what is a model?

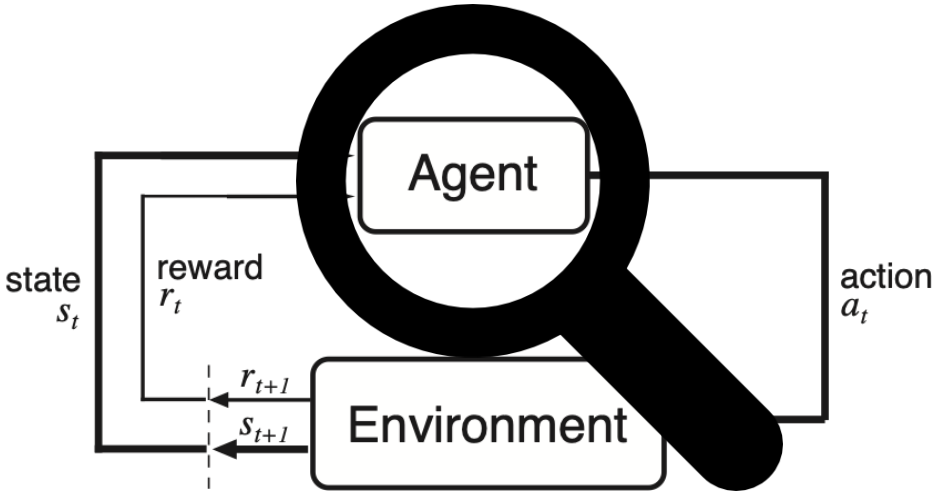

But what is a model?

Given a state and an action, a model returns the next state and next reward

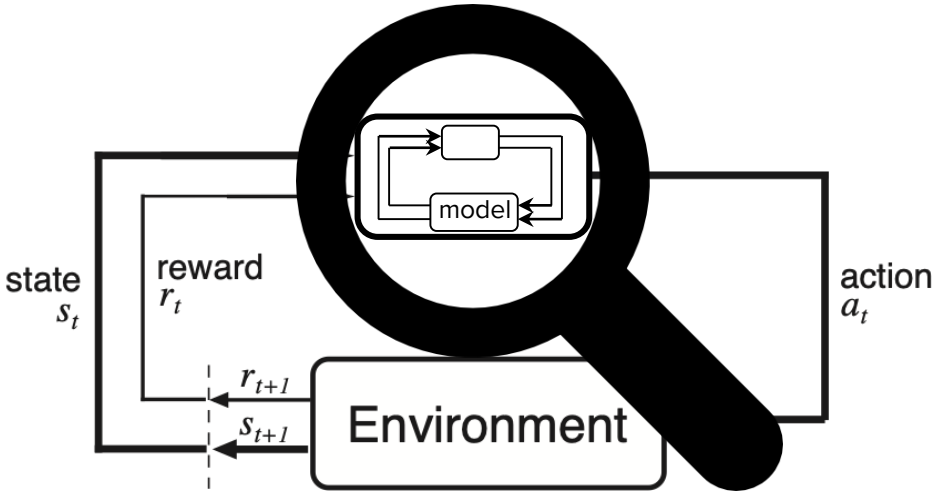

Model-free and Model-based RL

- Model-free RL:

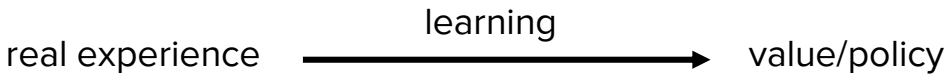

- Model-based RL:

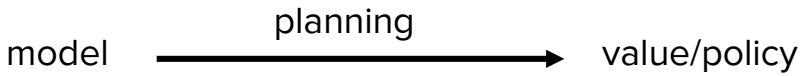

Integrating planning and learning

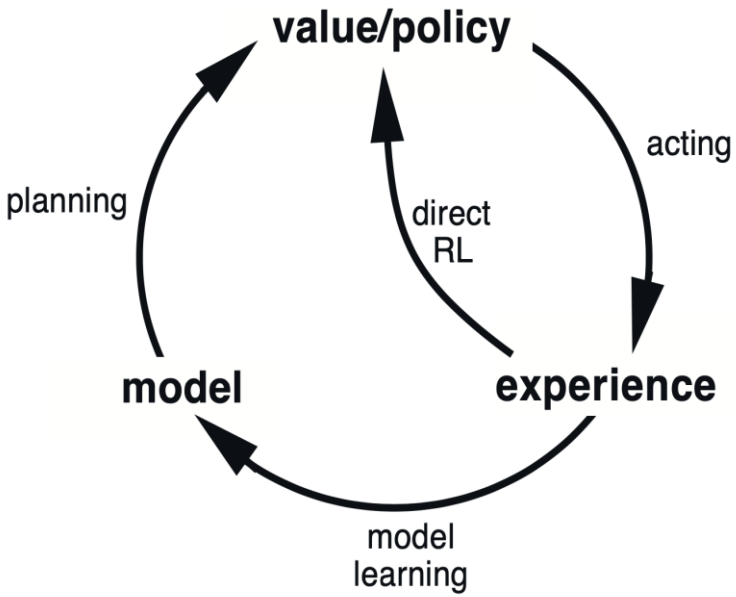

Dyna-Q architecture

Tabular Dyna-Q

Initialize $Q(s,a)$ and $Model(s,a)$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$

Loop forever:

(a) $S \leftarrow$ current (nonterminal) state

(b) $A \leftarrow \varepsilon$-greedy(S,Q)

(c) Take action A; observe resultant reward, R, and state, S'

(d) $Q(S,A) \leftarrow Q(S,A) + \alpha\big[R + \gamma \max_a Q(S',a) - Q(S,A)\big]$

(e) $Model(S,A) \leftarrow R, S'$ (assuming deterministic environment)

(f) Loop repeat n times:

$S \leftarrow$ random previously observed state

$A \leftarrow$ random action previously taken in S

$R, S' \leftarrow Model(S,A)$

$Q(S,A) \leftarrow Q(S,A) + \alpha\big[R + \gamma \max_a Q(S',a) - Q(S,A)\big]$

A mathematical framework for cognition?

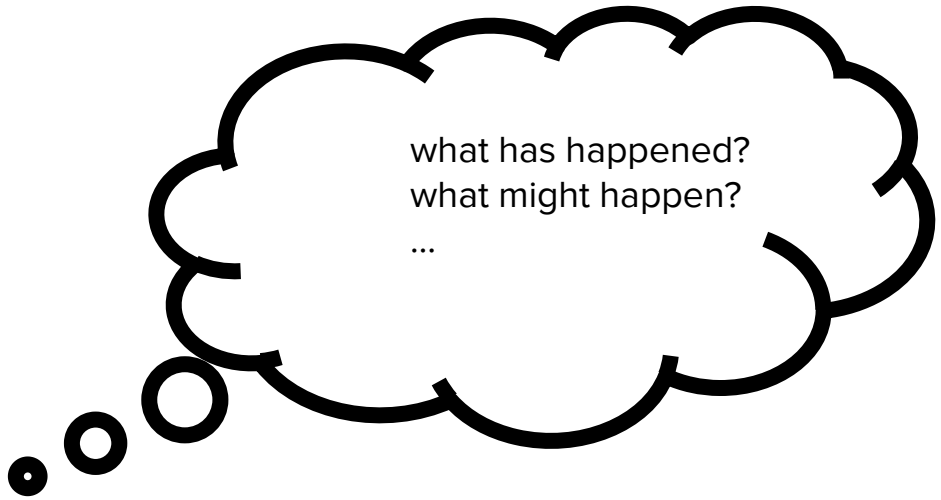

Replay for consolidation

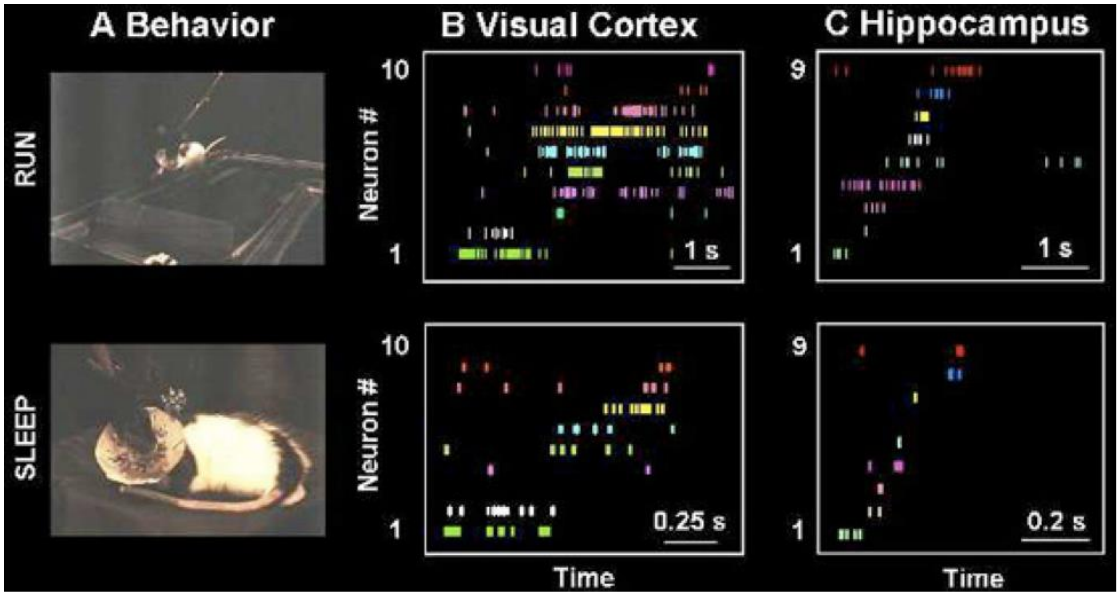

Ji, D. & Wilson, M.A. (2007)

Replay for planning

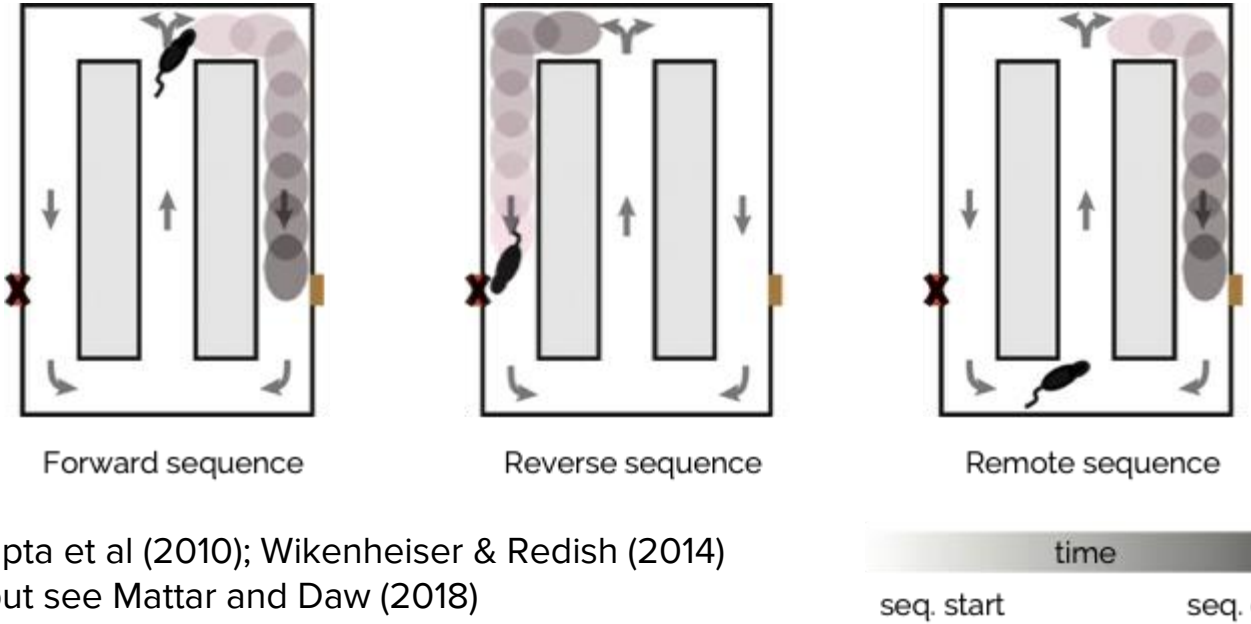

Gupta et al (2010); Wikenheiser & Redish (2014)
...but see Mattar and Daw (2018)